

setjmp, Longjmp and User-Level Threads

*Failures are much more fun to hear about afterward,
they are not so funny at the time.*

setjmp and longjmp

- ❑ You can jump between functions, conditionally!
- ❑ Use a jump buffer to save the return point.
- ❑ Use a long-jump to jump to a return point.
- ❑ Header file `setjmp.h` is required.
- ❑ A jump buffer is of type `jmp_buf`.
- ❑ Set up a jump buffer with function `setjmp()`.
- ❑ Execute a long jump with function `longjmp()`.

Here is the concept

- ❑ Declare a variable of type jmp_buf:
`jmp_buf JumpBuffer;`
- ❑ Call function setjmp () to mark a return point:
`setjmp (JumpBuffer) ;`
- ❑ Later on, use function longjmp () to jump back:
`longjmp (JumpBuffer, 1) ;`

The jump buffer is used in both calls

The meaning of this argument will be clear later

But, you need to know more!

- ❑ When `setjmp()` is called, it saves the current state of execution and returns 0.
- ❑ When `longjmp()` is called, it sends the control back to a marked return point and let `setjmp()` to return its *second* argument?????

```
#include <setjmp.h>
jmp_buf Buf;
void A(...)

{
    if (setjmp(Buf)==0) {
        printf("Marked!\n");
        /* other statements */
        B(...);      first time here
    }
    else {
        printf("Returned from"
               " a long journey\n");
        /* other statement */
    }
}

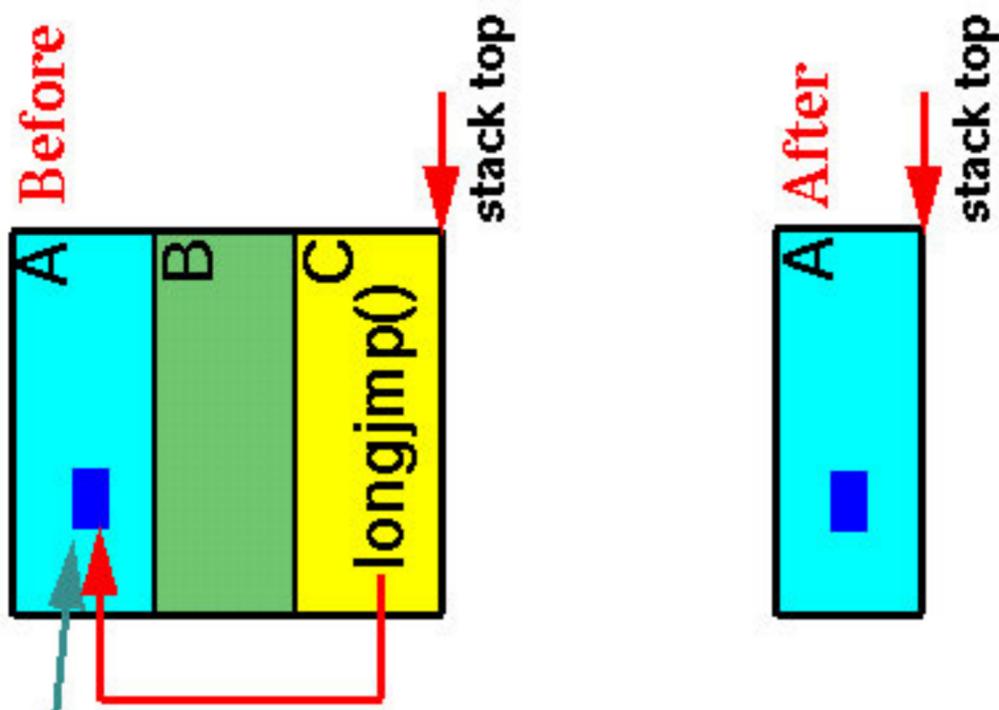
void B(...)
{
    /* other statement */
    longjmp(Buf, 1);
}
```

Control Flow of `setjmp()` and `longjmp()`

```
void A(...)  
{  
    if (setjmp(JumpBuffer) == 0)  
        B(...);  
    else {  
        printf("A long journey\n");  
        ....  
    }  
}
```

The content of a jmp buffer when execute a longjmp() must be valid

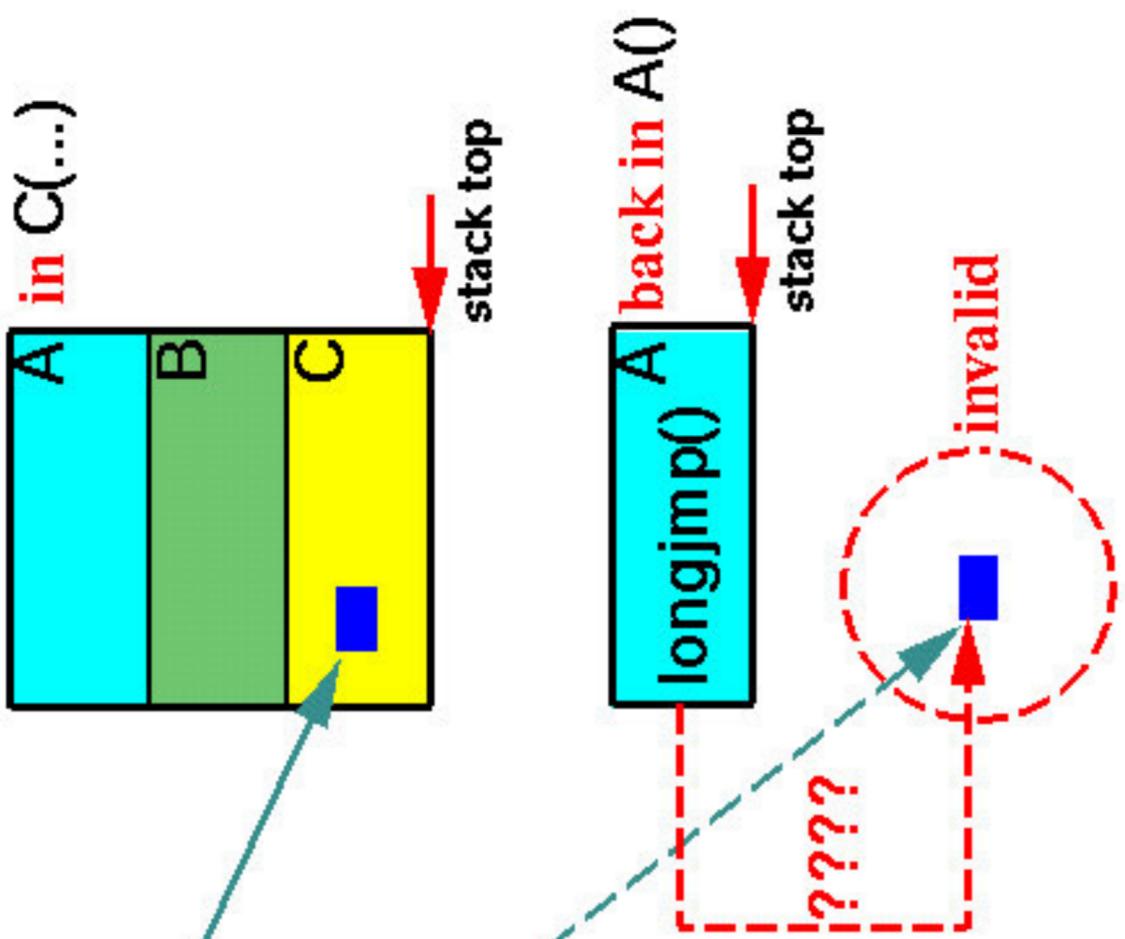
```
jmp_buf Buf;
void A(...){ if (setjmp(Buf) == 0)
    B(...);
else
    printf("I am back\n");
}
void B(...){ C(...) }
void C(...){ longjmp(Buf,1); }
```



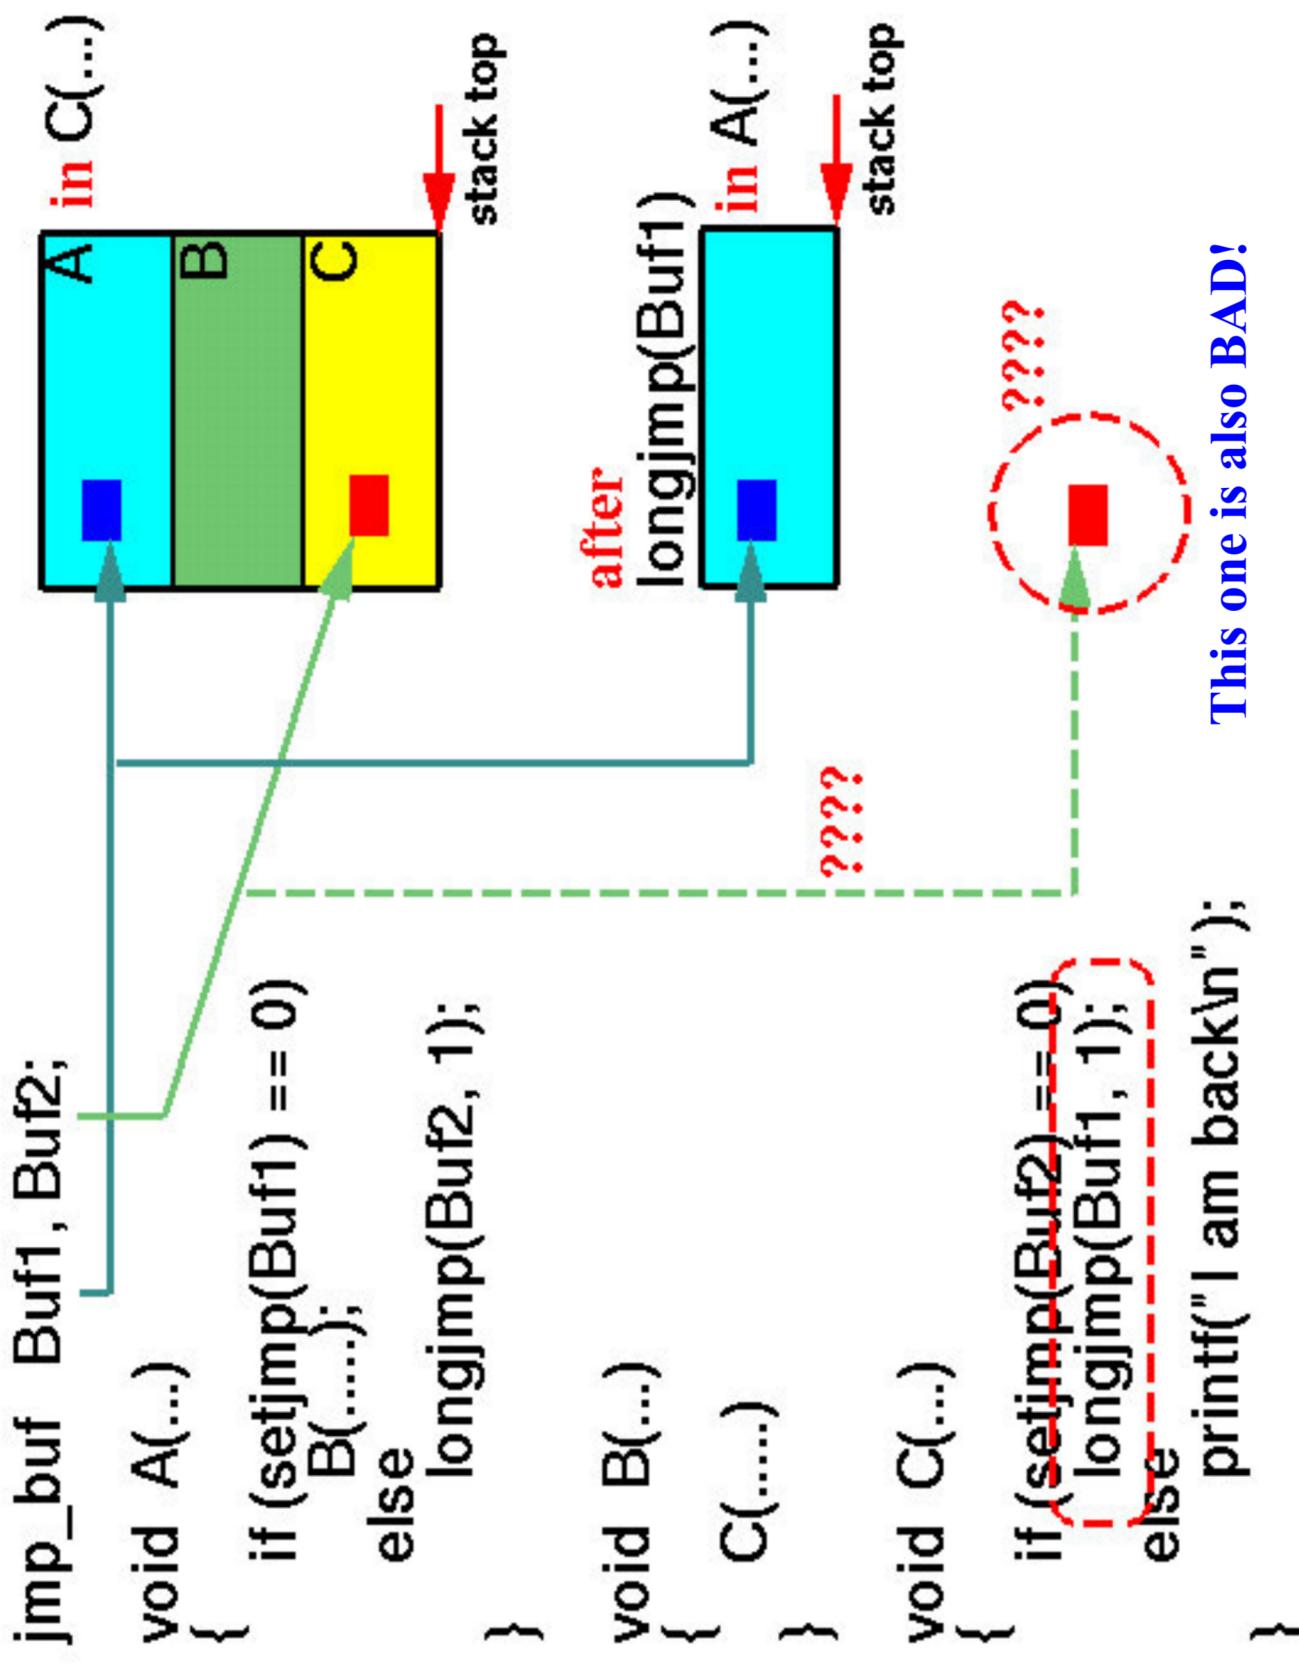
This is a good one!

The content of a jump buffer when execute a longjmp() must be valid

```
jmp_buf Buf;
void A(...){ B(...); longjmp(Buf,1); }
void B(...){ C(...); }
void C(...){
    if (setjmp(Buf) == 0)
        return;
    else
        printf("I am back\n");
}
```



The content of a jump buffer when execute a longjmp() must be valid



This one is also BAD!

Jump Buffer Example: Factorial: 1/3

```
#include <stdio.h>
#include <setjmp.h>

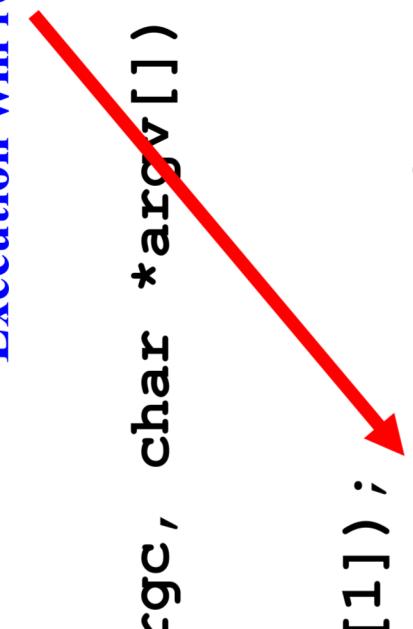
jmp_buf JumpBuffer;

int result;

void main(int argc, char *argv[])
{
    int n;
    n = atoi(argv[1]);
    if (setjmp(JumpBuffer) == 0)
        factorial(n);
    else
        printf("%d\n", n, result);
    exit(0);
}
```

Execution will return to here!

Result is in here!



Jump Buffer Example: Factorial: 2/3

```
void factorial(int n)
{
    fact(1, n);
}
```

```
void fact(int Result, int Count, int n)
{
    if (Count <= n)
        fact(Result*Count, Count+1);
    else {
        result = Result;
        longjmp(JumpBuffer, 1);
    }
}
```

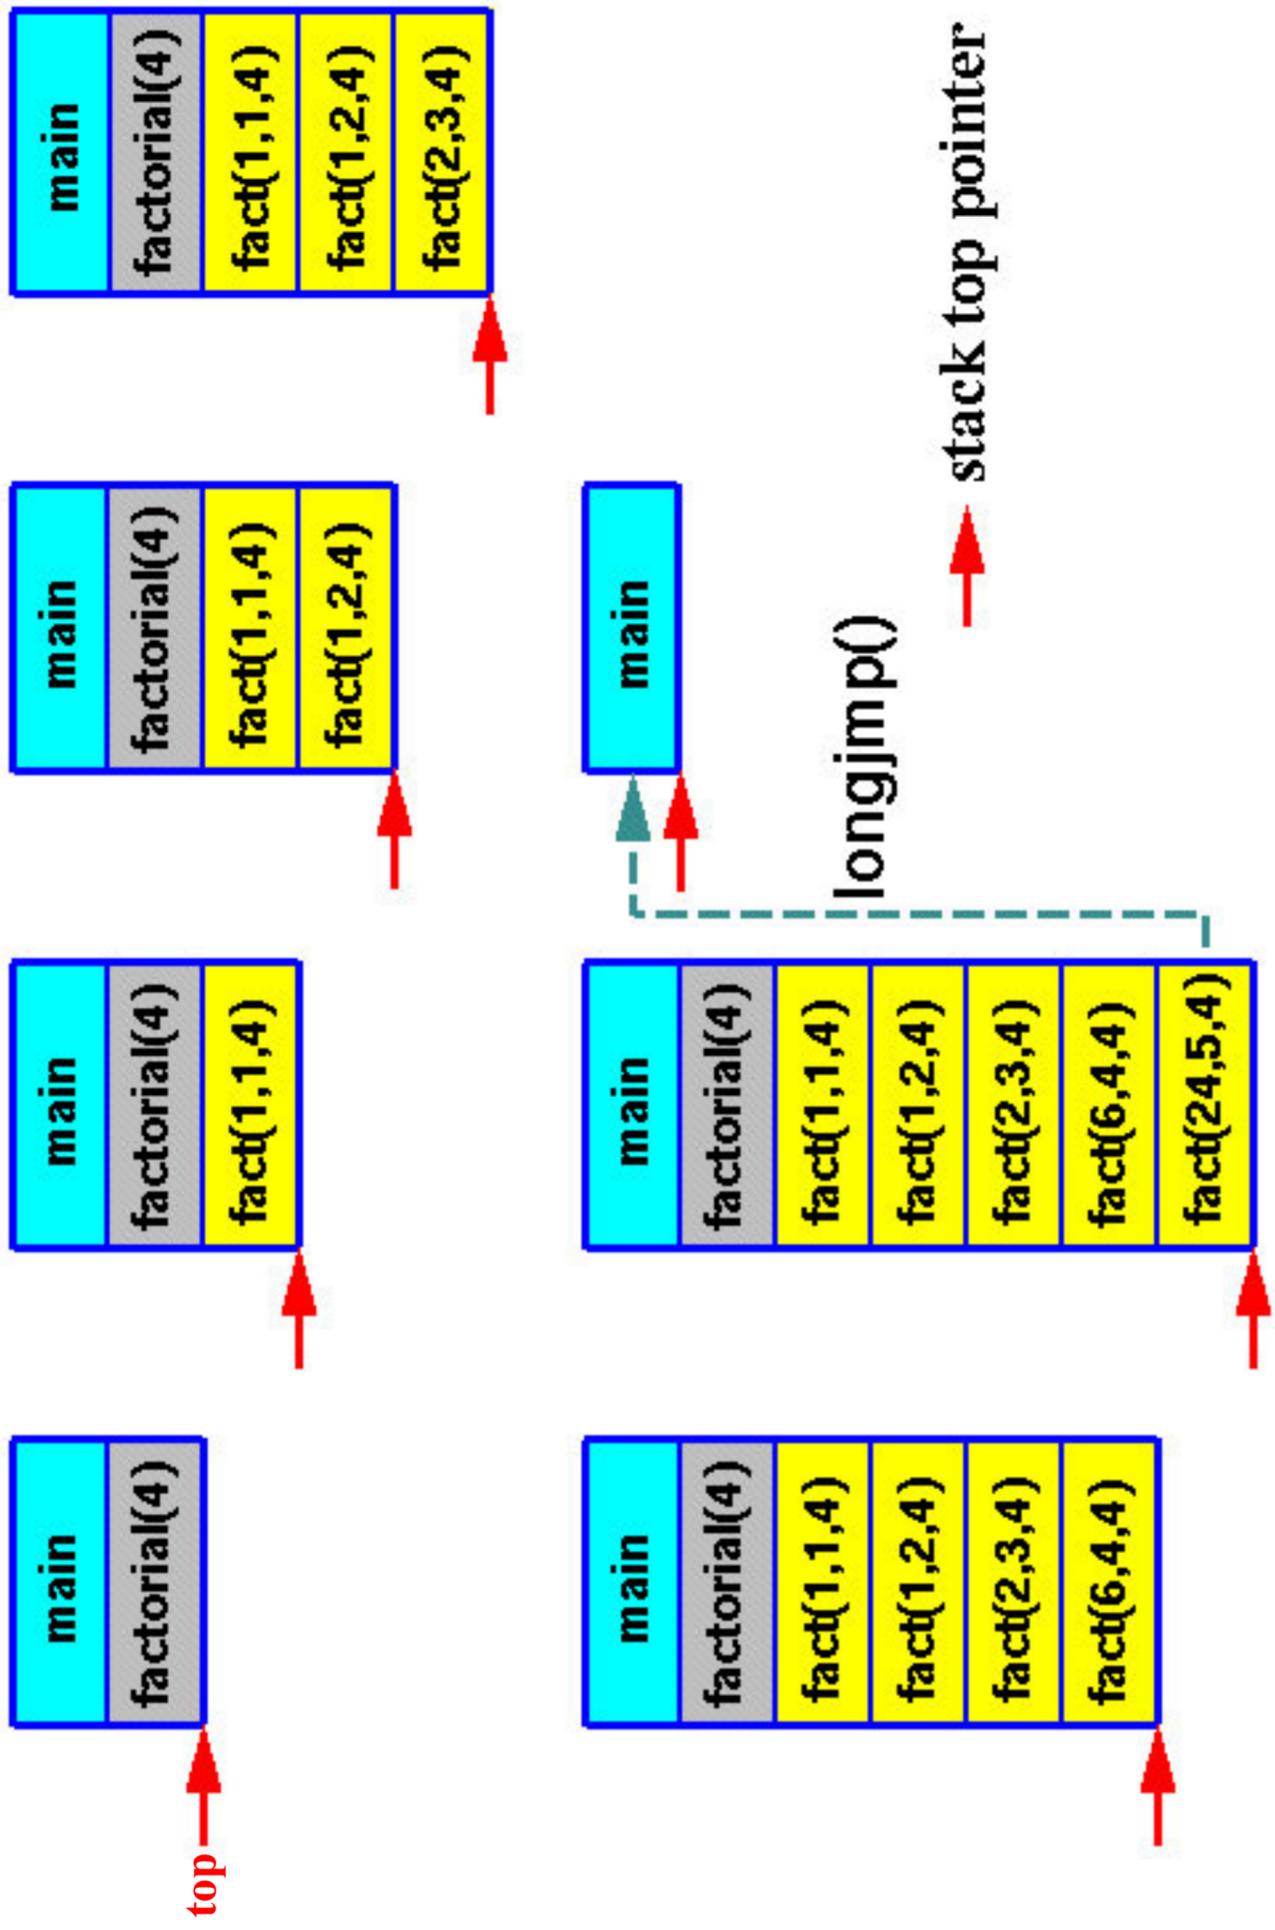
Why not Count++ or ++Count?

Count++: the value of current (rather than the next one) Count is passed
++Count: We don't know the evaluation order of the argument.

Left to Right: OK

Right to Left: Oops! The value of Count in Result*Count is wrong.

Jump Buffer Example: Factorial: 3/3



Jump Buffer Example: Signals-1

```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>

jmp_buf JumpBuffer;
void handler(int);

void main(void)
{
    signal(SIGINT, handler);
    while(1) {
        if (setjmp(JumpBuffer)==0) {
            printf("Hit Ctrl-C ...\\n");
            pause();
        }
    }
}
```

```
void handler(int sig)
{
    char c;
    signal(sig, SIG_IGN);
    printf("Ah, Ctrl-C?\\n");
    printf("Want to quit?");
    c = getchar();
    if (c == 'Y' || c == 'y')
        exit(0);
    else {
        signal(SIGINT, handler);
        longjmp(JumpBuffer,1);
    }
}
```

Without this `longjmp()`, control returns to here!

Jump Buffer Example: Signals-2

(1/2)

```
#define START 0
#define FROM_CTRL_C 1
#define FROM_ALARM 2
#define ALARM 5

jmp_buf Buf;
INT(int);
ALRM(int);
void
void main(void)
{
    int Return;
    signal(SIGINT, INT);
    signal(SIGALRM, ALRM);
}

while (1) {
    if ((Return=setjmp(Buf) )==START) {
        alarm(ALARM);
        pause();
    }
    else if (Return == FROM_CTRL_C) {
    }
    else if (Return == FROM_ALARM) {
        print("Alarm reset to %d sec.\n",
              ALARM);
        alarm(ALARM);
    }
}
```

Jump Buffer Example: Signals-2 (2/2)

```
void INT(int sig)
{
    char c;
    signal(SIGALRM, SIG_IGN);
    signal(SIGINT, SIG_IGN);
    signal(SIGALRM, ALRM);
    printf("Got an alarm\n");
    alarm(0); /* reset alarm */
    signal(SIGALRM, ALRM);
    signal(SIGINT, INT);
    longjmp(Buf FROM_ALARM);
}

void ALRM(int sig)
{
    signal(SIGINT, SIG_IGN);
    signal(SIGALRM, SIG_IGN);
    print("Want to quite?");
    c = getchar();
    if (c=='y' || c=='Y')
        exit(0);
    signal(SIGINT, INT);
    signal(SIGALRM, ALRM);
    longjmp(Buf, FROM_CTRL_C);
}
```

alarm clock has no effect

A Strange Use: 1/2

```
#include <stdio.h>
#include <setjmp.h>

int      max, iter;
jmp_buf  Main, PointA, PointB;
void    Ping(void), Pong(void);

void main(int argc, char *argv[])
{
    max = abs(atoi(argv[1]));
    iter = 1;
    if (setjmp(Main) == 0)
        Ping();
    if (setjmp(Main) == 0)
        Pong();
    longjmp(PointA, 1);
}
```

Set return point
Main & call Ping()

Set return point
Main & call Pong()

A Strange Use: 2/2

```
void Ping (void)
{
    if (setjmp (PointA)==0)
        longjmp (Main,1);
    while (1) {
        printf ("Ping--");
        if (setjmp (PointA)==0)
            longjmp (PointB,1);
    }
}
void Pong (void)
{
    if (setjmp (PointB)==0)
        longjmp (Main,1);
    while (1) {
        printf ("Pong--");
        if (iter++ > max)
            exit(0);
        if (setjmp (PointB)==0)
            longjmp (PointA,1);
    }
}
```

This program does not work if there are local variables. Why?

Output:

Ping-Pong-Ping-Pong-Pong-.....

A Not-So-Correct Multithread System: 1/10

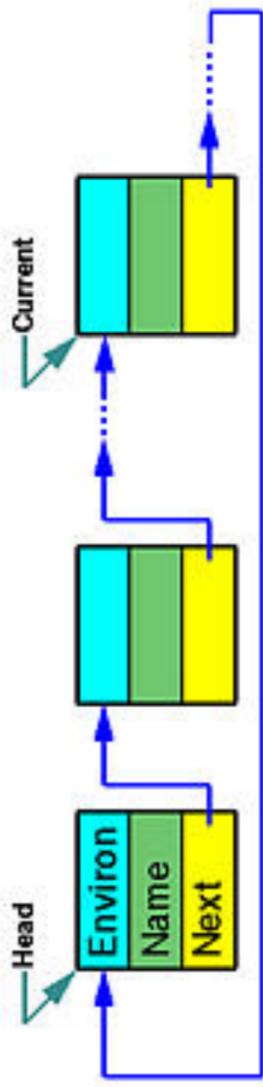
- ❑ Before going into more details, we examine a not-so-correct way to build a user-level thread system.
- ❑ First, we need a (simplified) TCB data structure.

```
typedef struct PCB_NODE *PCB_ptr; /* pointer to a PCB */  
  
/* a PCB:  
 * jmp_buf Environment;  
 * int Name;  
 * PCB_ptr Next;  
 */  
PCB;
```

A Not-So-Correct

Multithread System: 2/10

- ❑ We need two more jump buffers `MAIN` and `SCHEDULER`.
- ❑ The former is used to save the main program's environment, and the latter is for the scheduler.
- ❑ Because the main program and the scheduler are not scheduled by the scheduler, they are not in the PCB list.
- ❑ There are two pointers: `Head` pointing to the head of the PCB list and `Current` to the running thread.



A Not-So-Correct

Multithread System: 3/10

- The scheduler is simple.
- Initially the scheduler `Scheduler()` is called by the main program to set an entry in jump buffer `SCHEDULER` and jump back to the main program using jump buffer `MAIN` that was setup *before* the call to `Scheduler`.
- After this, we use a long jump to `SCHEDULER` rather than via function call.

```
void Scheduler(void)
{
    if (setjmp(SCHEDULER) == 0)
        longjmp(MAIN, 1);
    Current = Current->Next;
    longjmp(Current->Environment, 1); /* jump to its environ */
}
```

A Not-So-Correct

Multithread System: 4/10

- ❑ `THREAD_YIELD()` is very simple.
- ❑ Release CPU voluntarily.
- ❑ What we need is saving the current environment to this thread's environment (actually a jump buffer) and transferring the control to the scheduler via a `longjmp`.
- ❑ Because this is so simple, we use `#define`

```
#define THREAD_YIELD(name) {  
    if (setjmp(Current->Environment) == 0) {  
        longjmp(SCHEDULER, 1);  
    }  
}
```

A Not-So-Correct

Multithread System: 5/10

- ☐ `THREAD_INIT()` can be part of `THREAD_CREATE()`.
- ☐ We create and initialize a PCB, set its return point, and long jump back to the main program.

```
#define THREAD_INIT(name) {  
    work = (PCB_ptr) malloc(sizeof(PCB));  
    work->Name = name;  
    if (Head == NULL)  
        Head = work;  
    else  
        Current->Next = work;  
    work->Next = Head;  
    Current = work;  
    if (setjmp(work->Environment) == 0)  
        longjmp(MAIN, 1);  
}
```

A Not-So-Correct Multithread System: 6/10

- ☐ `THREAD_CREATE()` is simple.
- ☐ We just set the return point of `MAIN` and call the function.

```
#define THREAD_CREATE(function, name) { \
    if (setjmp(MAIN) == 0) \
        (function)(name); \
}
```

```
void main(void)
{
    Head = Current = NULL; /* initialize pointers */
    THREAD_CREATE(funct_1, 1); /* initialize threads */
    THREAD_CREATE(funct_2, 2);
    THREAD_CREATE(funct_3, 3);
    THREAD_CREATE(funct_4, 4);
    if (setjmp(MAIN) == 0) /* initialize scheduler */
        Scheduler();
    longjmp(SCHEDULER, 1); /* start scheduler */
}
```

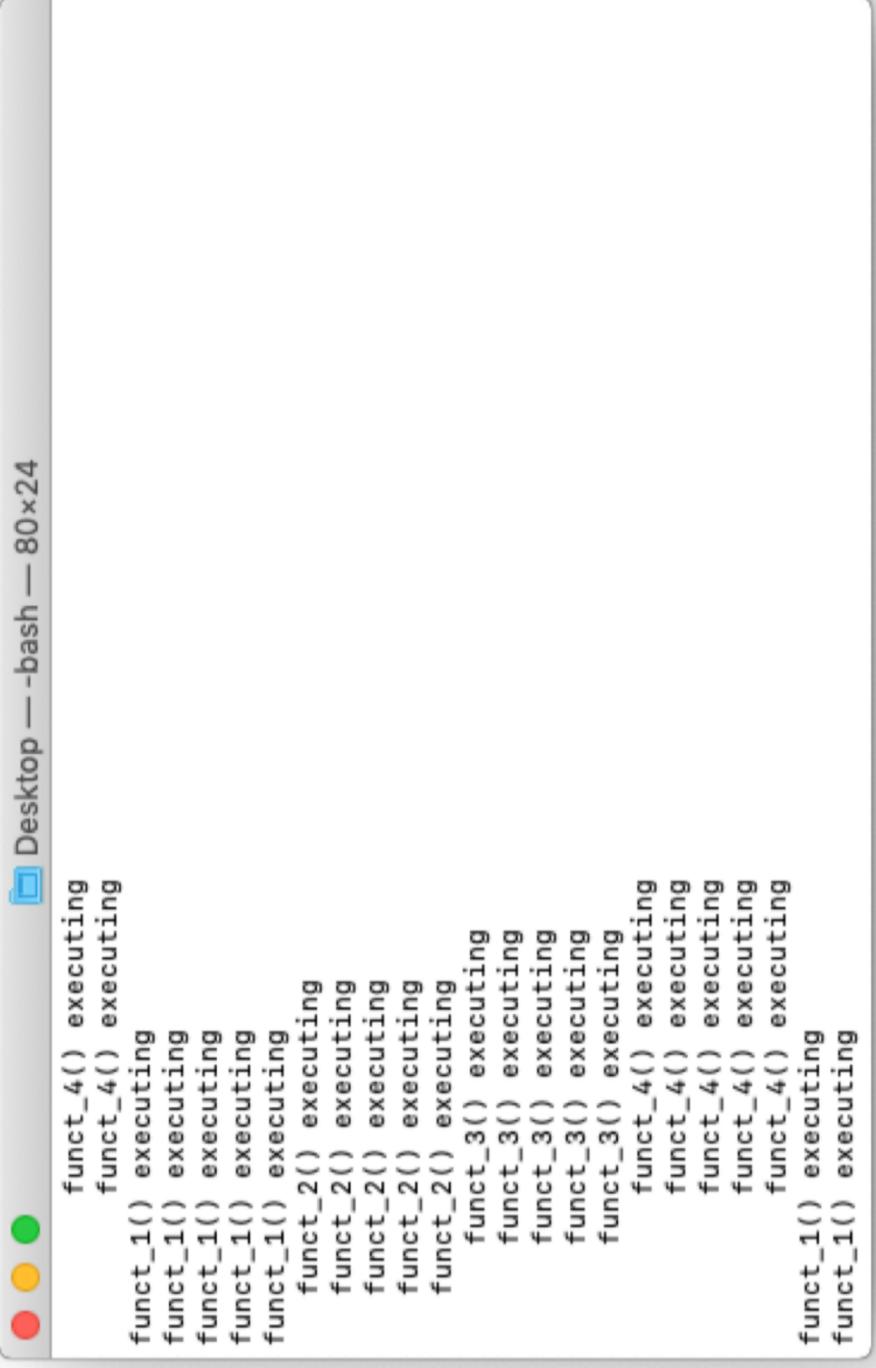
A Not-So-Correct Multithread System: 7/10

- Each function to be run as a thread must call `THREAD_INIT()`.

```
void funct_1(int name)
{
    int i;
    THREAD_INIT(name); /* initialize as thread */
    while (1) {          /* running the thread */
        for (i = 1; i <= MAX_COUNT; i++)
            printf("funct_1() executing\n");
        THREAD_YIELD(name); /* yield control */
    }
}
```

A Not-So-Correct Multithread System: 8/10

- ❑ This implementation appears to be correct. The following is a screenshot.



The screenshot shows a terminal window titled "Desktop — -bash — 80x24". The window contains the following text:

```
funct_4() executing
funct_4() executing
funct_1() executing
funct_1() executing
funct_1() executing
funct_1() executing
funct_1() executing
funct_1() executing
funct_2() executing
funct_2() executing
funct_2() executing
funct_2() executing
funct_2() executing
funct_2() executing
funct_3() executing
funct_3() executing
funct_3() executing
funct_3() executing
funct_3() executing
funct_3() executing
funct_4() executing
funct_4() executing
funct_4() executing
funct_4() executing
funct_4() executing
funct_1() executing
funct_1() executing
```

A Not-So-Correct

Multithread System: 9/10

- ❑ It is **not!** **Why?** But you have all the ideas!
- ❑ We do not use many local variables, in fact only one variable `i`. In a function, say `funct_1()`, `i` is used before `THREAD_YIELD()`.
- ❑ Once `THREAD_YIELD()` is called, the stack frame of `funct_1()` becomes invalid. However, this is fine, because after returning from `THREAD_YIELD()` this variable is reinitialized.
- ❑ Is the jump buffer Environment of each thread correct? In general it is not. However, the PC is correct because it is not stored there.

A Not-So-Correct Multithread System: 10/10

- ❑ The key issue making this system not-so-correct is that each thread does not have its stack frame.
- ❑ As a result, once it long jumps out of its environment the stack frame allocated by the system becomes invalid.
- ❑ The solution is simple: allocating a separate stack frame for each “created” thread so that it won’t go away.
- ❑ This is what we intend to do.

Let Us Solve the Problem: 1/10

- We need a better TCB. Env is a redefined jmpbuf.

```
typedef struct TCB_NODE *TCB_ptr;
typedef TCB_ptr THREAD_t;

typedef struct TCB_NODE { /* thread control block
    int     Name;          /* thread ID
    int     Status;         /* thread state
    Env    Environment;   /* processor context area
    void   *StackBottom;  /* bottom of stack attached
    int     Size;          /* stack size
    void   (*Entry) (int, char**); /* entry point (function)
    int     Argc;          /* # of arguments
    char   **Argv;         /* argument list
    Queue  JoinList;      /* joining list of threads
} TCB;
```

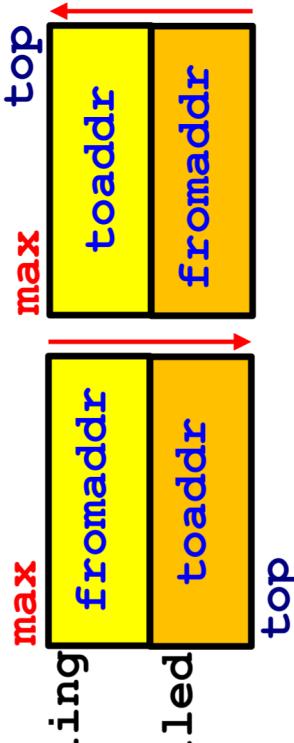
Initialize the Coroutines: 2/10

- ❑ Initialize the coroutines structure. Refer to slides used in CS3331 Concurrent Computing for the concept of coroutines.

```
static int THREAD_SYS_INIT (void)
{
    int *stack;
    Running = (THREAD_t) malloc(sizeof(TCB)); /* dummy running thread */
    stack = (int *) malloc(64); /* 64 bytes for stack */
    Running->Name = mtu_MAIN;
    Running->StackBottom = stack;
    ReadyQ_Head = ReadyQ_Tail = NULL;
    SuspendQ_Head = SuspendQ_Tail = NULL;
    SysTable_Head = SysTable_Tail = NULL;
    SYSTEM_INITIALIZE = TRUE;
    return mtu_NORMAL;
}
```

Stack Growing Direction: 3/10

- ❑ Because the user stack can grow up (stack at the highest address) or grow down (stack at the end of code and data sections), we need to know which way it goes.
- ❑ **fromaddr** is a variable in the calling function, and **toaddr** is a variable in the called function.



```
static int growsdown (void *fromaddr)
{
    int toaddr;

    return fromaddr > (void *) &toaddr;
}
```

Wrap the Created Thread: 4/10

- ❑ `THREAD_WRAP()` wraps up the created thread and runs it as a function. The function to be run is indicated by `Running`.

```
static volatile void THREAD_WRAP (void)
{
    (*Running->Entry) (Running->Argc, Running->Argv) ; /* run thread */
    THREAD_EXIT() ;
    /* if user did not call
     * THREAD_EXIT, do it here */
}
```

THREAD_INIT: System Dependent

- ❑ Initialize a new thread's environment.
- ❑ Newer version of gcc may not allow you to modify the jump buffer.

```
void THREAD_INIT (volatile TCB *volatile NewThread, void *StackPointer)
{
    /* In Linux 1.0 the code maybe like following three lines
     * NewThread->Environment->__sp = StackPointer;
     * NewThread->Environment->__bp = StackPointer;
     * NewThread->Environment->__pc = (void *)THREAD_WRAP;
     * Here is THREAD_INIT for Linux 2.0
     */
    NewThread->Environment[0].__jmpbuf[JB_SP] = (int)StackPointer;
    NewThread->Environment[0].__jmpbuf[JB_BP] = (int)StackPointer;
    NewThread->Environment[0].__jmpbuf[JB_PC] = (int)THREAD_WRAP;
}
```

THREAD_CREATE() : 1/2

- ❑ THREAD_CREATE() allocates a TCB and a stack for the thread being created and initialize the TCB.

```
THREAD_t THREAD_CREATE (void (*Entry) (), int Size, int Flag,
int Argc, char **Argv)
{
    THREAD_t NewThread;
    int *StackBottom, FromAddr;
    void *StackPointer;

    NewThread = (THREAD_t) malloc (sizeof(TCB)); /* new thread TCB */
    if (NewThread == NULL)
        return (THREAD_t) mtu_ERROR;
    Size += sizeof(StackAlign); /* get new stack size
StackBottom = (int *) malloc (Size);
StackPointer = (void *) (growsdown (&FromAddr) ?
    (Size+(int) StackBottom)-sizeof(StackAlign) : (int) StackBottom);
THREAD_INIT (NewThread, StackPointer); /* initialize thread
                                         32*/ /* architecture-dependent! */
```

[-- Next Page --](#)

THREAD_CREATE () : 2/2

- ❑ Continue from previous page.

```
/* from previous page */
NewThread->Name = NextThreadName++; /* initial TCB values */
NewThread->Status = mtu_READY;
NewThread->Entry = (void (*) (int, char**)) Entry;
NewThread->Argc = Argc; NewThread->Argv = Argv;
NewThread->StackBottom = StackBottom;
NewThread->Size = Size;
NewThread->JoinList.Head = NULL; NewThread->JoinList.Tail = NULL;
THREAD_READY (NewThread); /* thread into Ready Q */
THREAD_READY (Running);
if (Flag == THREAD_SUSPENDED)
    THREAD_SUSPEND (NewThread);
    THREAD_SCHEDULER ();
return NewThread;
}
```

THREAD_EXIT() : 1/2

- ❑ Continue from previous page.

```
int THREAD_EXIT (void)
{
    THREAD_t temp;

    if (Running->Name == mtu_MAIN) { /* if main, exit there
                                         /* have no thread remain
                                         /* in the ready Q */
        if (ReadyQ.Head != NULL)
            return mtu_ERROR;
        if (SuspendQ.Head != NULL) /* and in suspend queue
                                         /* and in suspend queue */
            return mtu_ERROR;
        return mtu_NORMAL;
    }

    while (Running->JoinList.Head != NULL) { /* check for joining
                                         /* temp = (THREAD_t) THREAD_Remove (&(Running->JoinList));
                                         /* temp->Status = mtu_READY;
                                         /* THREAD_Append (&ReadyQ, (void *) temp);
                                         /* continue to next page */
    }
}
```

THREAD_EXIT() : 2/2

- ❑ Continue from previous page.

```
Running->Name = mtu_INVALID;           /* set current thread's TCB */
Running->Status = mtu_TERMINATED;      /* and status = terminated */
Running = NULL;
THREAD_SCHEDULER();
return mtu_ERROR;
}
```

THREAD_YIELD()

- ❑ THREAD_YIELD() puts the running thread back to READY and calls THREAD_SCHEDULE() to reschedule.

```
void THREAD_YIELD (void)
{
    THREAD_READY(Running) ; /* put the running one to Ready */
    THREAD_SCHEDULE() ; /* ask scheduler to reschedule. */
}
```

THREAD_SCHEDULE () : 1/2

- ❑ THREAD_SCHEDULE () finds and runs the next ready thread.

```
static int THREAD_SCHEDULE (void)
{
    THREAD_t volatile Nextp;
    THREAD_t (THREAD_t) THREAD_Remove (&ReadyQ); /* find a thread */
    if (Nextp == NULL) {
        mtuMP_errno=mtuMP_DEADLOCK; /* if ready queue is empty */
        ShowDeadlock ();
        exit (0);
        return mtu_ERROR;
    }
    if (Running==NULL) {
        Running = Nextp;
        Nextp->Status = mtu_RUNNING;
        RestoreEnvironment (Running->Environment); /* restore env */
    }
    /* continue to next page */
}
```

THREAD_SCHEDULE() : 2/2

- ❑ THREAD_SCHEDULE() finds and runs the next ready thread.

```
/* continue from previous page */
if ((Running != NextTp) &&
(SaveEnvironment(Running->Environment) == 0)) {
    /* else save running's env */
    /* let next thread run */
    NextTp->Status = mtu_RUNNING;
    RestoreEnvironment(Running->Environment); /* restore env */
}
return mtu_NORMAL;
```

Dining Philosophers: 1/2

```
void Philosopher(int No)
{
    int Left = No;
    int Right = (No + 1) % PHILOSOPHERS;
    int RandomTimes, i, j;
    char spaces[PHILOSOPHERS*2+1];

    for (i = 0; i < 2*No; i++) /* build leading spaces */
        spaces[i] = ' ';
    spaces[i] = '\0';

    printf("%sPhilosopher %d starts\n", spaces, No);
    for (i = 0; i < Iteration; i++) {
        printf("%sPhilosopher %d is thinking\n", spaces, No);
        SimulatedDelay();
        SEMAPHORE_WAIT(Seats);
        printf("%sPhilosopher %d has a seat\n", spaces, No);
        MUTEX_LOCK(Chopstick[Left]);
        MUTEX_LOCK(Chopstick[Right]);
        printf("%sPhilosopher %d gets chopsticks and eats\n", spaces, No);
        SimulatedDelay();
        printf("%sPhilosopher %d finishes eating\n", spaces, No);
        MUTEX_UNLOCK(Chopstick[Left]);
        MUTEX_UNLOCK(Chopstick[Right]);
        SEMAPHORE_SIGNAL(Seats);
    }
    THREAD_EXIT();
}
```

Dining Philosophers: 2/2

```
int main(int argc, char *argv[])
{
    THREAD_t Philosophers[PHILOSOPHERS];
    int SeatNo[PHILOSOPHERS];
    int i;

    Iteration = abs(atoi(argv[1]));
    srand((unsigned int)time(NULL));
    /* initialize random number */
    for (i = 0; i < PHILOSOPHERS; i++) /* create mutex locks
    Chopstick[i] = MUTEX_INIT();
    Seats = SEMAPHORE_INIT(PHILOSOPHERS-1); /* seat semaphore */
    for (i = 0; i < PHILOSOPHERS; i++) { /* create philosophers */
        /* philosopher number */
        /* create a thread */
        /* the thread function */
        /* stack size */
        /* the thread flag */
        /* play a trick here */
        /* no argument list */
        if (Philosophers[i] == (THREAD_t)mtu_ERROR) { /* if failed */
            printf("Thread creation failed.\n");
        }
    }
    for (i=0; i<PHILOSOPHERS; i++) /* wait until all done */
        THREAD_JOIN(Philosophers[i]);
}

return 0;
}
```

Conclusions

- ❑ In a kernel, the kernel has to allocate a stack differently.
- ❑ Context switch has to be done differently and directly rather than using a jump buffers.
- ❑ The remaining should be very similar and could be copied easily.
- ❑ The not-so-correct system is discussed here:
<http://www.cs1.mtu.edu/cs4411.ck/www/NOTES/non-local-goto/index.html>
- ❑ A simple user-level thread system is in the common directory mtuthread.tar.gz. Also refer to this page:
<http://www.cs1.mtu.edu/cs4411.ck/www/PROG/PJ/proj.html>

The End